


# Training Deep Neural Networks

# What we have learnt so far?


---

- ▶ Training set is  $\{(\mathbf{x}^1, y^1), (\mathbf{x}^2, y^2), (\mathbf{x}^3, y^3) \dots (\mathbf{x}^m, y^m)\}$
- ▶  $L$  = number of layers in the network
- ▶  $s_l$  = number of units in layer  $l$



# Gradient descent steps

*How much does a small change in weights affect the empirical loss ?*



## Algorithm


1. Initialize weights randomly  $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3.     Compute gradient,  $\frac{\partial J(\theta)}{\partial \theta}$
4.     Update weights,  $\theta = \theta - \alpha \frac{\partial J(\theta)}{\partial \theta}$
5. Return weights

$$\frac{\partial J(\theta)}{\partial \theta^{(2)}} = \underbrace{\frac{\partial J(\theta)}{\partial y}}_{\text{purple}} * \underbrace{\frac{\partial z_1}{\partial \theta^{(2)}}}_{\text{red}}$$

$$\frac{\partial J(\theta)}{\partial \theta^{(1)}} = \underbrace{\frac{\partial J(\theta)}{\partial y}}_{\text{purple}} * \underbrace{\frac{\partial y}{\partial z_1}}_{\text{red}} * \underbrace{\frac{\partial z_1}{\partial \theta^{(1)}}}_{\text{blue}}$$

# What have we learnt so far?

Forward pass



$$\begin{aligned} a^{(1)} &= x \\ z^{(2)} &= \Theta^{(1)} a^{(1)} \\ a^{(2)} &= g(z^{(2)}) \left( \text{add } a_0^{(2)} \right) \\ z^{(3)} &= \Theta^{(2)} a^{(2)} \\ a^{(3)} &= g(z^{(3)}) \left( \text{add } a_0^{(3)} \right) \\ z^{(4)} &= \Theta^{(3)} a^{(3)} \\ a^{(4)} &= h_{\Theta}(x) = g(z^{(4)}) \end{aligned}$$

# What have we learnt so far? Backpropagation

---

- Compute  $\delta^4 = a^4 - y$ , then compute


$$\delta^{(3)} = (\Theta^{(3)})^T \delta^{(4)} \cdot * g'(z^{(3)})$$

$$\delta^{(2)} = (\Theta^{(2)})^T \delta^{(3)} \cdot * g'(z^{(2)})$$

- $\Theta^3$  is the vector of parameters for the layer 3
- $g'(z^3)$  is the first derivative of the activation function  $g$
- $g'(z^3) = a^3 \cdot * (1 - a^3)$  (derivative of sigmoid)
- $\delta^3 = (\Theta^3)^T \delta^4 \cdot *(a^3 \cdot * (1 - a^3))$
- $\delta^2 = (\Theta^2)^T \delta^3 \cdot *(a^2 \cdot * (1 - a^2))$

# Back propagation intuition

---




# Activation Functions

---

- ▶ What is an activation function?
  - ▶ an activation function is used to turn  $\mathbf{z}$  which is a linear value, into  $\mathbf{a}$ , a non-linear value.
- ▶ What happens without non-linear activation functions?
- ▶ Each layer can use different activation functions (in theory)

# Sigmoid Activation Function

- ▶ Sigmoid is normally used for final layer
- ▶ Intermediate layers can have other activation functions




$$a = g(z) = \frac{1}{1+e^{-z}}$$

$$g'(z) = g(z)(1 - g(z)) = a(1 - a)$$



# tanh


- ▶ The tanh function is very similar to the sigmoid function. It is actually just a scaled version of the sigmoid function.
- ▶ Tanh works similar to the sigmoid function but is symmetric over the origin. it ranges from -1 to 1.



$$a = g(z) = \tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$
$$g'(z) = 1 - (\tanh(z))^2 = 1 - a^2$$

# ReLU (Rectified Linear Unit)

- ▶ sigmoid and tanh functions  
gradient becomes very small if  
the value of  $z$  is positively large  
or negatively small, making the  
descent slow.
- ▶ The main advantage of using  
the ReLU function over other  
activation functions is that it  
does not activate all the  
neurons at the same time.




$$g(z) = \max(0, z)$$

$$g'(z) = \begin{cases} 0 & \text{if } z < 0 \\ 1 & \text{if } z \geq 0 \end{cases}$$

# Leaky ReLU

- ▶ But ReLU also falls a prey to the gradients moving towards zero.
- ▶ This can create dead neurons which never get activated
- ▶ Leaky ReLU is defined to address this problem.



$$g'(z) = \begin{cases} 0.01 & \text{if } z < 0 \\ 1 & \text{if } z \geq 0 \end{cases}$$

# Softmax function

---

- The softmax function is also a type of sigmoid function but is handy when we are trying to handle classification problems.

$$\sigma(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}} \quad \text{for } j = 1, \dots, K.$$

# Which activations functions to use when?

---

- ▶ Sigmoid functions and their combinations generally work better in the case of classifiers
- ▶ Sigmoids and tanh functions are avoided due to the vanishing gradient problem
- ▶ ReLU function is a general activation function and is used mostly
- ▶ If you encounter dead neurons in our networks the leaky ReLU function is the best choice
- ▶ ReLU function should only be used in the hidden layers
- ▶ As a rule of thumb, you can begin with using ReLU function and then move over to other activation functions in case ReLU doesn't provide with optimal results

# Gradient Checking

---

- ▶ Gradient checking is used to double-check that our gradient calculation is correct.

$$d\theta_{approx} = \frac{J(\theta_i + \epsilon) - J(\theta_i - \epsilon)}{2\epsilon}$$

- ▶ Use it sparingly as it is very slow.

- ▶ For each feature  $i$ , calculate the approximate gradient

$$diff = \frac{\|d\theta_{approx} - d\theta\|_2}{\|d\theta_{approx}\|_2 + \|d\theta\|_2}$$

- ▶ Then calculate the difference with the actual gradient calculated during gradient descent:

# More Optimization Algorithms

# Gradient Descent with Momentum

---

- ▶ With this method, we update  $\theta$  with the exponentially weighted running average of  $d\theta$
- ▶ This will work faster than the standard gradient descent.
- ▶ The intuition is that by averaging the values, it smoothens the oscillation in the gradient descent
- ▶ Here  $V_{d\theta}$  is an exponentially weighted moving average of  $d\theta$ . It's initialized to 0 and updated on every iteration:

$$V_{d\theta} = \beta V_{d\theta} + (1 - \beta)d\theta \quad \theta = \theta - \alpha V_{d\theta}$$



# RMSprop (Root Mean Square Prop)

---

- The intuition is to dampen large oscillation in some particular directions by penalizing movements that are large (dividing it by square root of itself). The formulas are:

$$S_{d\theta} = \beta S_{d\theta} + (1 - \beta)(d\theta)^2$$

$$S_{db} = \beta S_{db} + (1 - \beta)(db)^2$$

$$\theta = \theta - \alpha \frac{d\theta}{\sqrt{S_{d\theta} + \epsilon}}$$

$$b = b - \alpha \frac{db}{\sqrt{S_{db} + \epsilon}}$$

- where  $\beta$  is typically 0.999 and epsilon  $\epsilon$  is very small (e.g.  $10^{-8}$ ) and is used to avoid division by zero.

# Adam

- ▶ Adam (adaptive moment estimation) optimization combines momentum and RMSprop optimizations above.

- ▶ This is the momentum part with bias correction:

$$V_{d\theta} = \beta_1 V_{d\theta} + (1 - \beta_1) d\theta \quad S_{d\theta} = \beta_2 S_{d\theta} + (1 - \beta_2) (d\theta)^2$$

$$V_{d\theta}^{corr} = \frac{V_{d\theta}}{1 - \beta_1^t} \quad S_{d\theta}^{corr} = \frac{S_{d\theta}}{1 - \beta_2^t}$$

$$V_{db} = \beta_1 V_{db} + (1 - \beta_1) db \quad S_{db} = \beta_2 S_{db} + (1 - \beta_2) (db)^2$$

$$V_{db}^{corr} = \frac{V_{db}}{1 - \beta_1^t} \quad S_{db}^{corr} = \frac{S_{db}}{1 - \beta_2^t}$$

$$\theta = \theta - \alpha \frac{V_{dV}^{corr}}{\sqrt{S_{d\theta}^{corr} + \epsilon}} \quad b = b - \alpha \frac{V_{bcr}}{\sqrt{S_{db}^{corr} + \epsilon}}$$

- ▶ The typical hyperparameter values are  $\beta_1 = 0.9$ ,  $\beta_2 = 0.999$ , learning rate  $\alpha$  needs to be tuned, and epsilon  $\epsilon$  is  $10^{-8}$ .

# Which one to choose?

---

- ▶ <https://towardsdatascience.com/types-of-optimization-algorithms-used-in-neural-networks-and-ways-to-optimize-gradient-95ae5d39529f>

# Learning Rate Decay

---

- ▶ Learning rate decay can be used to speed up learning
- ▶ Start with higher learning rate and reduce proportional to epoch number
- ▶ it can be calculated as follows, with  $\alpha_0=0.2$  and  $decay\_rate=1$ :

$$\alpha = \frac{1}{1+decay\_rate*epoch\_num} \alpha_0$$

# Batch Normalization

---

- ▶ The idea is to normalize the value of  $Z$  in any hidden layers
- ▶ This makes the values not fluctuate much between iterations
- ▶ This makes the subsequent layers better

# Bias and Variance

---

- ▶ When our training set and test set have different distribution, we may have the following results:
  - ▶ training error: 1%
  - ▶ test error: 10%
  - ▶ Is this bias or variance?
  - ▶ How to fix it?
- ▶ When both training and testing error both are high
  - ▶ Is this bias or variance?
  - ▶ How it fix it?




# Regularization

## ► How is bias and variance effected by regularization?

### Linear regression with regularization

Model:  $h_{\theta}(x) = \theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \frac{\lambda}{2m} \sum_{j=1}^m \theta_j^2$$



# How to choose $\lambda$ ?

---

- ▶ Have a set or range of values to use, for e.g, increment by factors of 2 so 0.01, 0.02, 0.04..., 10
- ▶ For each Minimize  $\lambda$  the cost function using gradient descent
- ▶ So now we have a set of parameter vectors corresponding to models with different  $\lambda$  value
- ▶ Check the cost function value for cross validation set and choose the  $\lambda$  that minimizes cross validation error





# Learning curves

- ▶ Useful to plot for algorithmic sanity checking or improving performance
- ▶ What is a learning curve?

$$J_{train}(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

